

Machine Learning Engineer

Assessment Instructions:

- Objective

- Use data/train_test.csv as the labeled development data.
- Decide how to split and validate the development data.
- Explore, clean, engineer features, validate, and model the data as you see fit.
- Run your trained model on every load in data/validation.csv.
 - Data/validation.csv contains the 12,000 loads requiring final predictions. Each row contains the load features and a unique load_id.
 - Run your trained model on every row in data/validation.csv.
 - data/validation_predictions_template.csv contains the same 12,000 load_id values.
 - Fill the template's predicted_rate column and save the completed file as validation_predictions.csv.
- Please refer to the README.md for setup instructions.

- Submit

- An accessible GitHub repository containing your solution code, dependencies, and run instructions.
- validation_predictions.csv containing exactly:
 - load_id,predicted_rate
- A PDF or DOCX report containing:
 - Your approach to validate test/train data and split.
 - The fixed December prediction chart produced by the provided score.py
- A 2-3 minute loom going over the following:
 - Key findings from exploring the data
 - Any data-quality issues identified and how you addressed them
 - Reasoning behind the chosen model
 - Training and validation approach, including how the data was split
 - A brief walkthrough of the most important parts of your code